

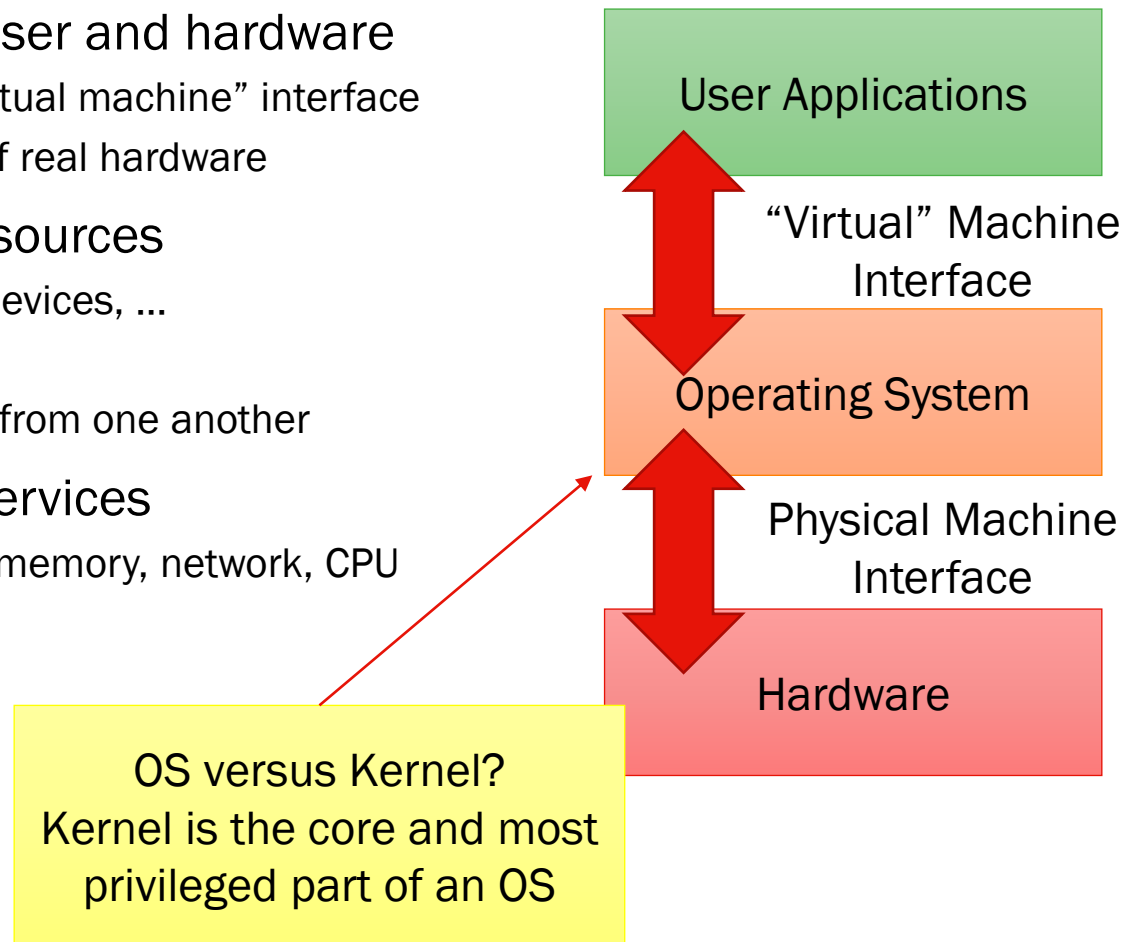
CS5460: Operating Systems

Lecture 1: Course Overview

(Chapter 2)

What is an Operating System?

- Interface between user and hardware
 - Exports simpler “virtual machine” interface
 - Hides ugly details of real hardware
- Manages shared resources
 - CPU, memory, I/O devices, ...
 - Ensure fairness
 - Protects processes from one another
- Provides common services
 - File system, virtual memory, network, CPU scheduling, ...
- Goals:
 - Convenient to use
 - Efficient



OS Techniques for Easier Apps

- **Abstraction**

- Often over physical resources: cores, memory, disk, net (**virtualization**)
- Sometimes others: pipes, process groups, ulimits, futexes

- **Advantages?**

- Portability
 - Applications run on more diverse hardware
- Higher-level
 - Easier to express complex applications
 - May be able to accelerate implementations without changing applications

Make sure you can name the abstractions for each of these resources

- **Challenges**

- What are the right abstractions? (often **virtualizations**)
- What is the cost of abstraction?

OS Techniques for Easier Apps

- Resource management and sharing
 - Inter application protection and isolation
 - Easy allocation and sharing of resources
 - Fair access to resources
- Challenges
 - What are the right **mechanisms**?
 - For safe isolation?
 - For sharing resources?
 - What are the right **policies**?
 - Efficiency, fairness?

What is the difference?

What are examples?

What is the benefit of separating them?

Our route: Three Easy Pieces

- **Virtualization**

- Processes, system calls, context switch, scheduling
- Address spaces, virtual addressing, virtual memory

- **Concurrency**

- Threads, locking primitives

- **Persistence**

- I/O Devices, Disks, Flash, File systems
- NFS and a bit of networking and distributed systems

Each of these topics first arose in the context of operating systems

Virtualizing Cores

```
$ numactl -C 1 /bin/bash
```

```
$ while true; do sleep 1; echo A; done & \  
  while true; do sleep 1; echo B; done &
```

```
[1] 27883
```

```
[2] 27884
```

```
A
```

```
B
```

```
A
```

```
B
```

```
A
```

```
B
```

Virtualizing Memory

```
void main() {  
    int* p = malloc(sizeof(int));  
    while (1) {  
        sleep(1);  
        printf("%p = %d\n", p, *p);  
        (*p)++;  
    }  
}
```

```
$ ./mem & ./mem &  
[1] 28481  
[2] 28482  
0x555555559260 = 0  
0x555555559260 = 0  
0x555555559260 = 1  
0x555555559260 = 1
```

Persistence

```
void main() {  
    int fd = open("foo",  
                  O_RDWR | O_CREAT | O_TRUNC,  
                  S_IRWXU);  
  
    assert(fd != -1);  
    int rc = write(fd, "test\n", 5);  
    assert(rc == 5);  
    close(fd);  
}
```

```
$ ./io  
$ cat foo  
test
```

No mention of devices,
disks, drivers,
hardware.

Where/how are I/O
devices accessed
here?

What keeps programs
from e.g. skipping the
filesystem code that
does permissions
checks?

Managing Concurrency

```
int i = 0;
```

```
void* run(void* _) {  
    for (int j = 0; j < 1000000; j++) i++;  
}
```

```
void main() {  
    pthread_t t1, t2;  
    pthread_create(&t1, NULL, run, NULL);  
    pthread_create(&t2, NULL, run, NULL);  
    pthread_join(t1, NULL);  
    pthread_join(t2, NULL);  
    printf("%d\n", i);  
}
```

```
$ ./inc  
1041048  
$ ./inc  
1087180
```



Please don't
write code like
this

Why Take OS?

- Remove doubt about how programs work
 - We take abstractions for granted, but they have implementations
 - Threads, processes, libraries, files, sockets
 - You should understand what happens when
 - A program is run
 - A process accesses a file
 - A process uses more memory than is physically available
 - A mouse is moved and the pointer moves on the screen
- Write better, more efficient code
 - Good code at one layer requires understanding how the layer below will carry things out
 - Learn performance engineering; OS is performance obsessed
- Learn interface design and tradeoffs by example
- Demystify kernel programming

This Course

- **Instructor:** Ryan Stutsman
- **TAs:** John, Amit, Vinita, and Calvin
- **Textbook:** Operating Systems: Three Easy Pieces v1.0
 - ostep.org
- **Required background:**
 - Undergrads: CS 4400 and ability to program in C
 - Grads: Ability to program in C
- Office hours on calendar
 - (Instructor Tue 2-3:30 & Thu 10:45-12:15)

Lecture 1 - Introduction (Ch 2)

Jan 19, 9:10am - 10:30am

Calendar CS 5460-001 Spring 2021 Operating Systems

Details

[Live Zoom Lecture Link](#)

Reading: [Chapter 2 - Introduction](#)

[Slides](#)

Delete

Edit

9:10a Lecture 1 - Introduction
(Ch 2)

Canvas

- Canvas will be used for announcements
 - Make sure you are setup to get notified of announcements via email
- Primarily **use Canvas Discussions** for questions
 - In general, try to keep discussions open
 - Send private messages when necessary
- Canvas Inbox if you need to contact us directly

Office Hours & MS Teams

- Office hours are listed near top of syllabus
- MS Teams "Office Hours" channel
 - TAs will create a thread when they are available
 - Message them there or directly during their hours
 - They can chat, screen share, etc.
- Will also use this answer questions during exam
- Use Canvas Discussions for most general Q&A

Lectures

- Live lecture via Zoom
 - Ask questions via Zoom chat
 - I won't be able to answer all questions; will do my best
 - I will have my camera on
 - Your camera is up to you; if you have it on be mindful
 - You can unmute, but respect the question flow in chat
- Lectures will be recorded
 - Posted with ≤ 24 hours delay on Canvas
- **Attend, participate, and pay attention**
 - Only required attendance: Midterm (3/16) & Final (5/5)

Assignments

- 2 to 4 Homeworks (10%)
 - Shorter and more focused on reinforcing concepts
- 4 or 5 Assignment (40%)
 - You'll write C code
 - **Some of them are time consuming – start early**
 - Graded on CADE lab Linux machines; get account if needed
<http://www.cade.utah.edu/>
- Exams (50%): Tie together concepts from lectures and projects
 - Students do well on projects; places extra importance on exams
 - Covers OSTEP readings, homework and projects, and lecture content
 - Open-book on Canvas during class/final exam time (3/16, 5/5)

Grading & Late Policy

- Standard 90/80/70/60 grading scale
- Late turn in excused for family/medical emergency
 - Please make arrangements in advance, if possible
- Each student has five late days to use on programming assignments only (Assignment 1-5)
 - Notify Lead TA for assignment how many late days you want to use for that assignment
- Late assignments incur 10% penalty per day (after late days have been applied)
 - ≤ 1 day late, 90% credit; ≤ 2 days, 80%; ≤ 3 days, 70%
 - > 3 days late receives 60% credit

Collaboration vs. Cheating

- To participate in this course you must:
 - Read SoC Policy on Academic Misconduct (see syllabus)
 - **Have an Acknowledgement Form filed in SoC office by end of second week otherwise you must withdraw or receive an EU**
- Do not...
 - Copy code from another student
 - Even look at code from another student
 - Copy code from the web
 - Ask for answers on StackOverflow or a similar web site
- It is okay to discuss solution strategies with classmates
- All forms of cheating/misconduct result in course failure
- We will check for cheating

Start early on assignments!

Classroom Expectations


- Zoom, Canvas, and Teams are our classroom
 - Pay one another respect you would in any classroom
- Examples of potential harassment/misconduct
 - Sending unwanted and unsolicited messages
 - "Pinning" video of students other than instructor or speaker
 - Unflattering, intimidating, or demeaning statements about anyone especially remarks related to race, color, origin, religion, sex, gender identity/expression, sexual orientation, background, veteran status, educational status, genetics

Concerns? Tell me if you feel comfortable doing so
Else: dept. advisors or Office of Equal Opportunity

A Brief History of Operating Systems

Prehistory (pre-1945)


- Charles Babbage (1792-1871) & Ada, Countess of Lovelace (1815-1852)
 - 1st computer architect/programmers
 - First digital computer
 - “Analytical engine”
 - Never actually got it to work (although others subsequently have)
 - No operating system: programmer programmed to raw hardware



Babbage Analytical Engine

1930s, 1940s

- Slow human “computers” → faster machines
- Mechanical relays, vacuum tubes, plug-boards, core memory:
 - Turing (the “Bombe”)
 - Aiken (Harvard architecture)
 - Von Neumann (Princeton IAS)
 - Eckert and Mauchley (ENIAC)
 - Zuse (Z1, Z3)
- Huge, hot, fragile, and slow by modern standards



Von Neumann

OS History: Phase I

Hardware is very expensive, humans are cheap!

- One user at the console
 - One function at a time (no overlap between computation and IO)
 - User sitting at console to debug
 - First OSES: Common library routines

Rent \$20,000 a month *in 1953*
~\$180k/mo today

IBM Model 701 (Early 1950's)



16 KHz, 2048 words of 36 bits, ~75 pixels

Unprogrammed

P1


Wait for Human Input

P2

OS History: Phase I

- Batch processing: load, run, print, dump, repeat
 - Put up glass walls around computer
 - Users give program (cards or tape) to human who schedules jobs
 - OS loads, runs, and dumps user jobs
 - Non-interactive batch processing
 - Efficient use of HW, at least while not loading decks
 - Debugging hard, core dumps
 - Short jobs starve behind large ones
- No overlap of I/O and compute


Uniprogrammed




Lost interactive
debugging →
Humans spend
more time waiting

OS History: Phase I

- Data channels and interrupts (early ‘DMA’)
 - Buffering and interrupt handling in OS (“batch monitor”)
 - No protection – one job running at a time!
 - Overlaps running computation and IO in parallel
 - “Spooling”
 - Users carried around permanent storage (cards, tapes, ...)



Unprogrammed
but overlapped I/O



IBM 7094 (Early 1960's)

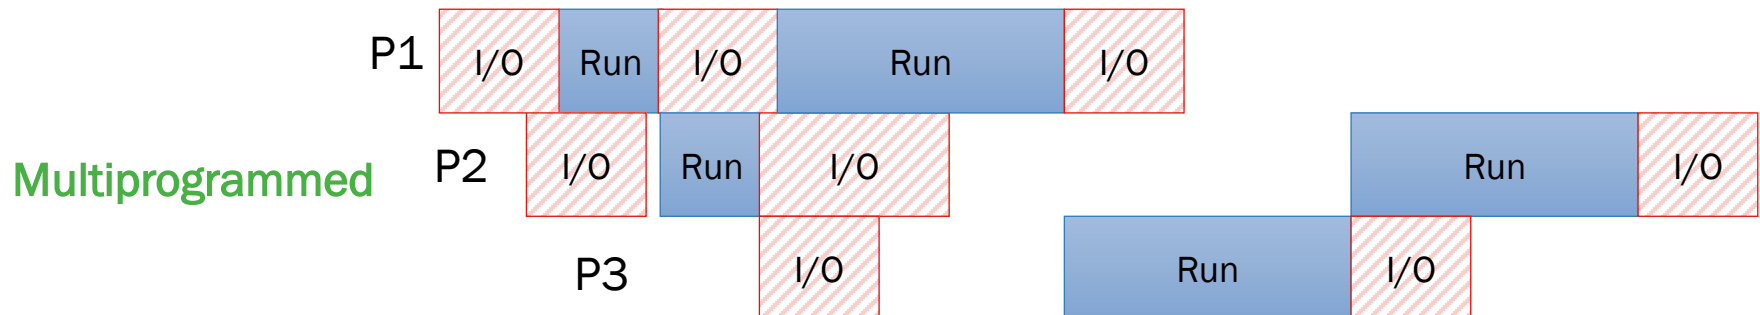


Just
\$400k/mo!

OS History: Phase I



- **Multiprogramming**; memory protection
 - More memory – can load several jobs at once
 - OS (monitor) always resident to coordinate activities
 - OS manages interactions between concurrent jobs:
 - Runs programs until they block due to I/O
 - Decides which blocked jobs to resume when CPU freed
 - Protects each job's memory from other jobs
 - IBM/360 combined IBM 1401 and IBM 7094
 - First machine to use ICs instead of individual transistors
- OS disasters → software engineering



OS History: Phase II (1965-1980)

Hardware is cheap, humans are expensive!

- Timesharing
 - First time-share system: CTSS from MIT (1962)
 - Timer interrupts: enable OS to take control (pre-emptive multitasking)
 - MIT/Bell Labs/GE collaboration led to MULTICS
 - Envisioned one huge machine for all of Boston (!!!)
 - Started in 1963, “done” in 1969, dead shortly thereafter
 - Bell Labs bailed on project, GE bailed on computers!
 - DEC PDP minicomputers: start of bottom feeding frenzy
 - PDP-1 in 1961 (4K 18-bit words, \$120,000)
 - Thompson and Ritchie
 - “C” language developed for Unix
 - Guiding principle of UNIX: Keep it simple so it can be built

OS History: Phase II (1965-1980)

- Timesharing (continued)
 - **Terminals are cheap**
 - Let all users interact with the system at once
 - **Debugging gets a lot easier**
 - Process switching occurs much more frequently
 - New OS services:
 - Shell to accept interactive commands, debuggers
 - File system to store data persistently
 - Virtual memory to allow multiple programs to be resident
 - New problems: response time and thrashing
 - Need to limit number of simultaneous processes or you can fall off performance cliff (“login”)
 - Lots of work on resource scheduling (CPU, memory, I/O)

OS History: Phase III (1980-2000s)

- Personal computing: every “terminal” has computer
 - One user per machine (remind you of anything?)
 - Initial PC OSes similar to old batch systems (w/TSR hacks)
 - Advanced OS features crept back in!
 - Linux, macOS, and Windows (starting with NT) now all include the ideas pioneered in the earlier decades



Original IBM PC



A young Bill Gates

OS History: Phase IV (2000s—now)

- Lots and lots of computers per person
 - Embedded systems
 - Cars commonly have 50+ processors
 - Cars, airplanes, factories run a huge amount of software
 - Mobile computing
 - PCs exceed the needs of many current computer users
 - Rise of smart phones and tablets
 - Cloud computing and serverless
 - Virtualized compute resources flexibly allocated on demand
 - Computing as a service rather than devices

The Future of OSes

- The very small
- The very large

	1983	2016	
MIPS	0.5	200,000+	400,000x
\$/MIP	\$100,000	\$0.005	20,000,000x
Memory	1 MB	3+ TB	3,000,000x
Network	0.1 Mbps	100,000 Mbps	1,000,000x
Storage	20 MB	10s of TB	500,000x
Address size	16 bits	64 bits	

- Change
 - 1953-2003: 10 orders of magnitude
 - 1983-2016: →
 - Nothing like it in other fields:
 - Transportation: 100x
 - Communication: 10^7

Characteristics of Modern OSes

- Enormous
 - Millions of lines of code, 1000s of engineer years
- Complex
 - Asynchronous, hardware idiosyncrasies, performance hacks
- Poorly understood
 - Systems outlive their builders
 - Never debugged: OS/360 released with 1000 bugs
 - Behavior hard to predict, security flaws
 - Unreliable
- Incredibly successful
 - Hard to imagine modern programming without them

Next time: Processes

- OSTEP Chapters 4, 5, 6

Important Terms and Ideas

- Operating System and Kernel
- Abstraction and Virtualization
- Policy and Mechanism
- Uniprogramming and Multiprogramming