

CS5460: Operating Systems

Lecture 3: Context Switch

(Chapters 4, 5, 6)

Slide Credit: Andrea Arpaci-Dusseau

Assignment 1

- Due Tue Feb 2
 - 1 week left!

Providing Good CPU Performance?

Direct execution

- Allow user process to run directly on hardware
- OS creates process, transfers control to start point (i.e., main())

Problems with direct execution?

1. Process could do something restricted
 Could read/write other process data (disk or memory)
2. Process could run forever (slow, buggy, or malicious)
 OS needs to be able to switch between processes
3. Process could do something slow (like I/O)
 OS wants to use resources efficiently and switch CPU to other process

Solution:

Limited direct execution – OS & hw maintain some control

Problem #1: Restricted Ops

How can we ensure user process can't harm others?


Solution: **privilege levels** supported by hw (status bit)

- User processes run in user mode (restricted mode) (Ring 3)
- OS runs in kernel mode (not restricted) (Ring 0)
 - Instructions for interacting with devices
 - Access to all memory
 - Ability to reconfigure CPU control registers (IDT, PTBR/CR3)

How can processes access devices?


- System calls (function call implemented by OS)
- Change privilege level through system call (trap)

System Call




P wants to call read()

System Call



P can only see its own memory because of user mode (other areas, including kernel, are hidden)


System Call



P wants to call read() but no way to call it directly


~~`callq sys_read`~~

System Call




```
movl $6, %eax;    int $64
```


System Call



System Call



`movl $6, %eax;`


`int $64`

syscall-table index

trap-table index

Trap instruction → kernel mode, vectors to trap handler
Kernel mode: we can do anything!

System Call



`movl $6, %eax;`


`int $64`

syscall-table index

trap-table index

Follow entries to correct system call code

System Call



`movl $6, %eax;`

`int $64`

syscall-table index

trap-table index

Kernel can access user memory to fill in user buffer
return-from-trap at end to return to Process P

What do we need to limit?

User processes are not allowed to perform:

- General memory access
- Disk I/O
- Special x86 instructions like `lidt`

What if process tries to do something restricted?

Problem #2: Take CPU Away?

OS requirements for **multiprogramming**
(multitasking):

- **Policy**: Decision-maker optimizing a performance metric
 - Process **Scheduler**: Which process when?
- **Mechanism**: Low-level code that implements the decision
 - **Dispatcher** and **Context Switch**: How?

Example of separation of policy and mechanism

Dispatch Mechanism

OS runs **dispatch loop**

```
while (1) {  
    run process A for some time-slice  
    stop process A and save its context  
    load context of another process B  
}
```

Context-switch


Question 1: How does OS/dispatcher gain control?

Question 2: What execution context must be
saved /restored?

Q1: How does OS get control?

Option 1: Cooperative Multitasking

- Trust process to relinquish CPU to OS through traps
 - Examples: System call, page fault (access page not in main memory), or error (illegal instruction or divide by zero)
 - Provide special `yield()` system call



Q1: How does OS get control?

- Problem with cooperative approach?
- Disadvantages: Processes can misbehave
 - By avoiding all traps and performing no I/O, can take over entire machine
 - Only solution: Reboot!
- Not performed in modern operating systems

Q1: How does OS get control?

Option 2: Preemptive Multitasking

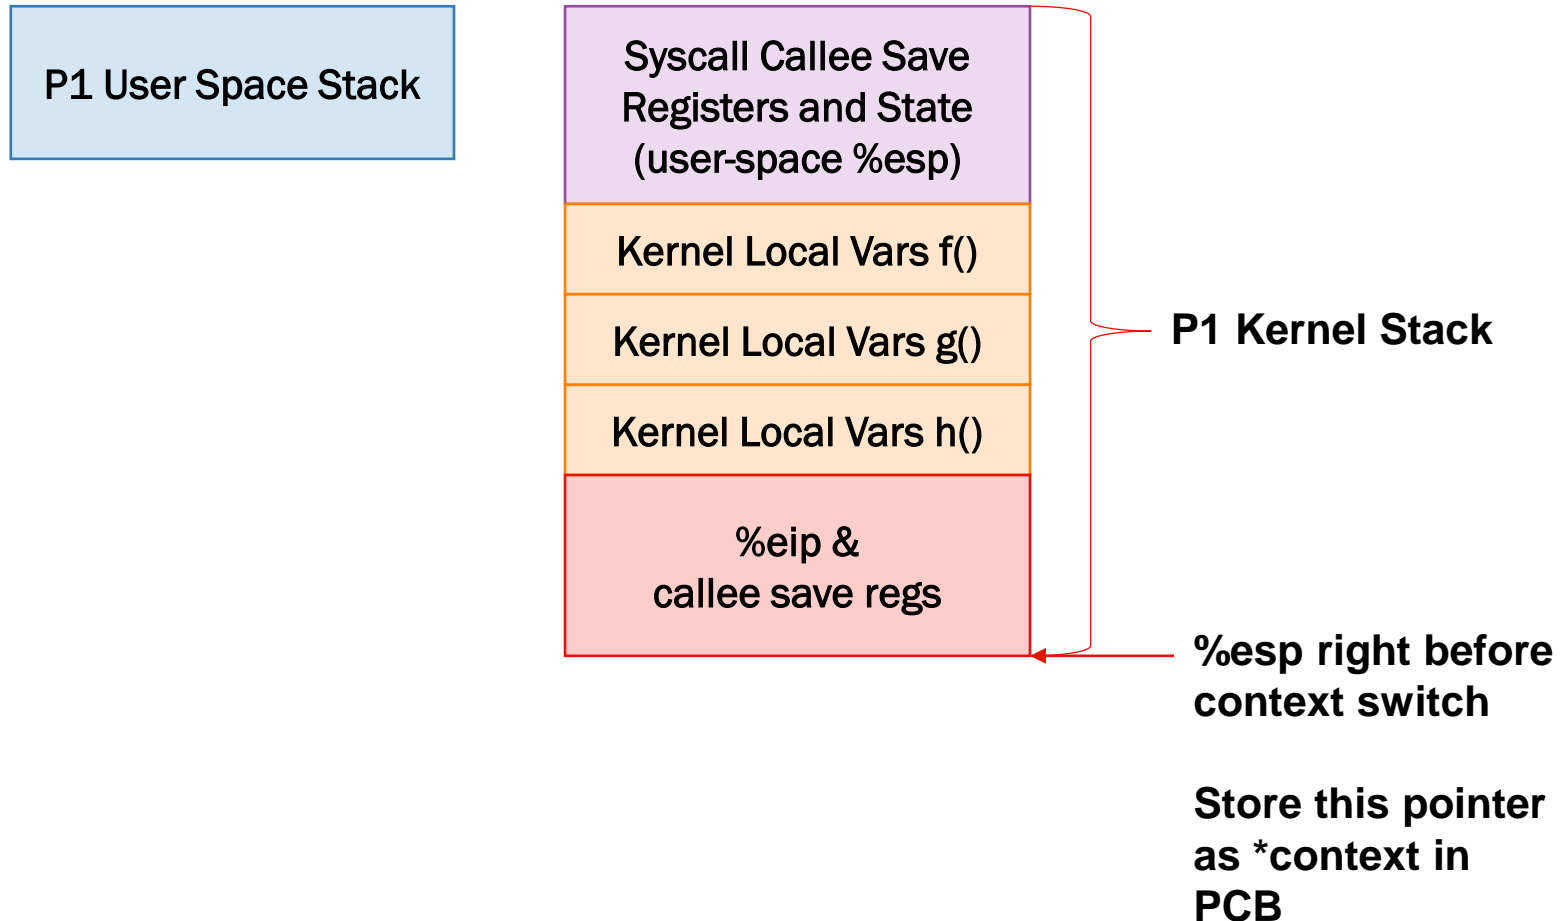
- Guarantee OS can obtain control periodically
- Enter OS by enabling periodic alarm clock
 - Hardware generates timer interrupt (CPU or separate chip)
 - Example: Every 10ms
- User must not be able to mask timer interrupt
- Dispatcher counts interrupts between context switches
 - Example: Waiting 20 timer ticks gives 200 ms time slice
 - Common time slices range from 4 ms to a few hundred ms

Q2: What context to save?

- **Process Control Block:** where dispatcher stores context of process when not running; contains
 - PID
 - Process state (i.e., running, ready, or blocked)
 - **Execution state (all registers, instruction ptr, stack ptr)**
 - Scheduling priority
 - Accounting information (parent and child processes)
 - Credentials (which resources can be accessed, owner)
 - Pointers to other allocated resources (e.g., open files)
- On fork: allocate PCB, initialize, put on ready queue (queue of runnable processes)
- On exit: clean up all process state (close files, release memory, page tables, etc)

How this stuff is handled is a bit tricky

Zooming In: Registers & KStack



Operating System

Hardware

Program

Process A

...

Syscall or timer interrupt
Hw switches to kstack
Raises to kernel mode
Save regs(A) to kstack(A)
Jump to trap handler

Handle the trap
Call **swtch()** routine
Save regs(A) to PCB(A)
Restore regs(B) from PCB(B)
Switch to kstack(B)
Return-from-trap (into B)

Restore regs(B)
from kstack(B)
Move to user mode
Jump to B's IP

Process B

...

xv6 PCB

```
enum proc_state { UNUSED, EMBRYO, SLEEPING,
                  RUNNABLE, RUNNING, ZOMBIE };

struct proc {
    uint sz; // Proc mem size (bytes)
    pde_t* pgdir; // Page table
    char* kstack; // Bottom of kstack
    enum procstate state; // Process state
    int pid; // Process ID
    struct proc* parent; // Parent process
    struct trapframe* tf; // Trap frm for syscall
    struct context* context; // swtch() here to run
    void* chan; // If !0, sleep on chan
    int killed; // If !0, been killed
    struct file* ofile[NOFILE]; // Open files
    struct inode* cwd; // Current directory
    char name[16]; // Process name
};

struct context
{
    uint edi;
    uint esi;
    uint ebx;
    uint ebp;
    uint eip;
};
```

Context Switch

- Context switches are fairly expensive
 - Time sharing systems do 100-1000 context switches per second
 - When? Timer interrupt, packet arrives on network, disk I/O completes, user moves mouse, ...
- lab2-15 3.8 μ s
- gamow 1.6 μ s
- home 1.0 μ s
- How might one go about measuring this?

Problem #3: Slow Ops (I/O)?

On op that does not use CPU,
OS switches to other processes

OS must track process states:

Running:


On the CPU (1 on a uniprocessor)

Ready:

Waiting for the CPU


Blocked:

Asleep: Waiting for I/O or
synchronization to complete



Problem #3: Slow Ops (I/O)?

- OS maintains queues of all PCBs
 - **Ready queue**: Contains all ready processes
 - Event queue: One logical queue per event
 - e.g., disk I/O and locks
 - Contains all processes waiting for that event to complete
- Invariant: each process in 1 state and on 1 queue



CPU Virtualization Summary

- Virtualization:
Context switching gives each process impression it has its own CPU
- Direct execution makes processes fast
- Limited execution at key points to ensure OS retains control
- Hardware provides a lot of OS support
 - user vs kernel mode
 - timer interrupts
 - automatic register saving on syscall

Process Management

- OS manages processes:
 - Creates, deletes, suspends, and resumes processes
 - Schedules processes to manage CPU allocation
 - Manages inter-process communication and synchronization
 - Allocates resources to processes (and takes them away)
- Processes use OS functionality to cooperate
 - Signals, sockets, pipes, files to communicate

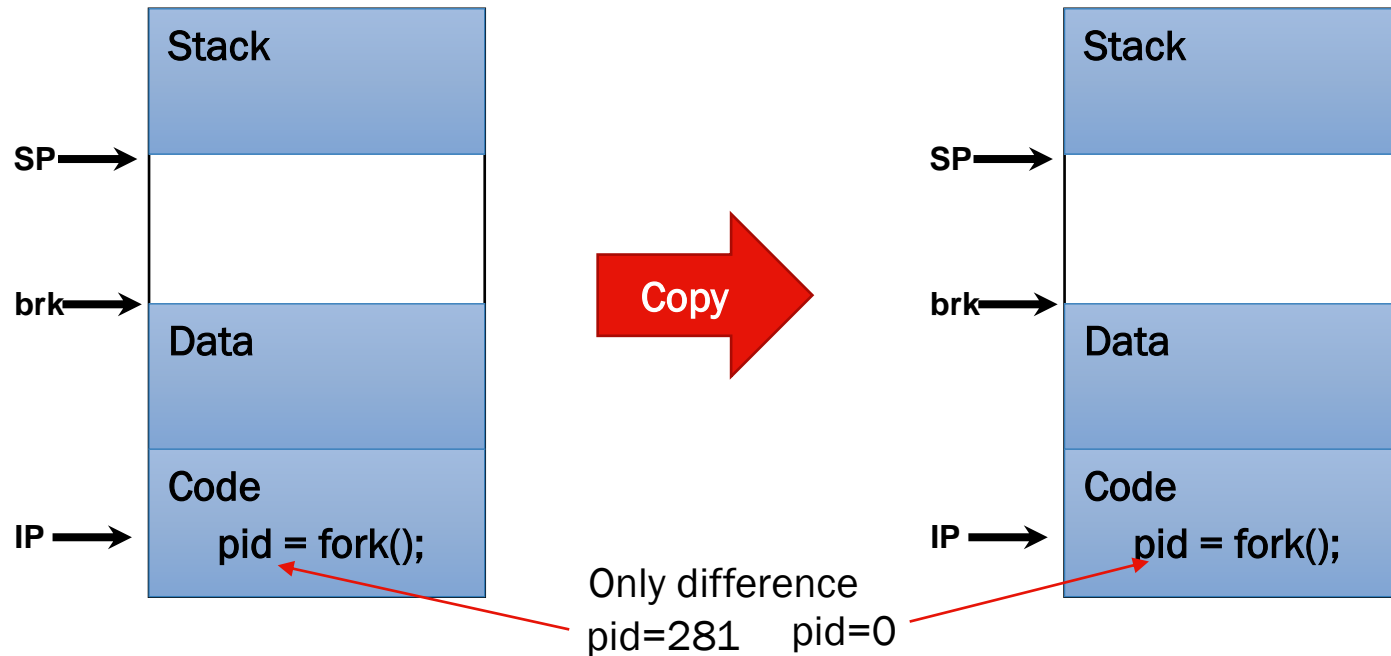
Practical Process Management

- On a Unix machine, try:
 - `ps -Af`: lots of info on all running processes
 - `kill -9 547`: terminates process with PID 547
 - `top` : displays dynamic info on top running jobs
 - Write a program that calls:
 - `getpid()`: returns current process's PID (process id)
 - `fork()`: create a new process
 - `wait()`: wait for exit of a child process
 - `exec()`: load a new program into the current process
 - `sleep()`: puts current process to sleep for specified time
- Commands work on macOS
- On Windows → Task manager (CTL-ALT-DEL)

Creating Processes: fork()

- Creates process that is near-clone of forking parent
Address space and running state is cloned
- Return of fork() differs:
0 in child
child PID in parent
- Many kernel resources are shared
open files and sockets
- wait() lets a process wait for the exit of a child
- To spawn new program, use some form of exec()

Semantics of fork()



- `fork()`, `exit()`, and `exec()` are weird!
 - `fork()` returns twice – once in each process
 - `exit()` does not return at all
 - `exec()` usually “does not return”: replaces process’ program

fork() and wait()

```
int main(int argc, char *argv) {  
    printf("parent %d\n", (int)getpid());  
    pid_t rc = fork();  
    assert(rc >= 0)
```

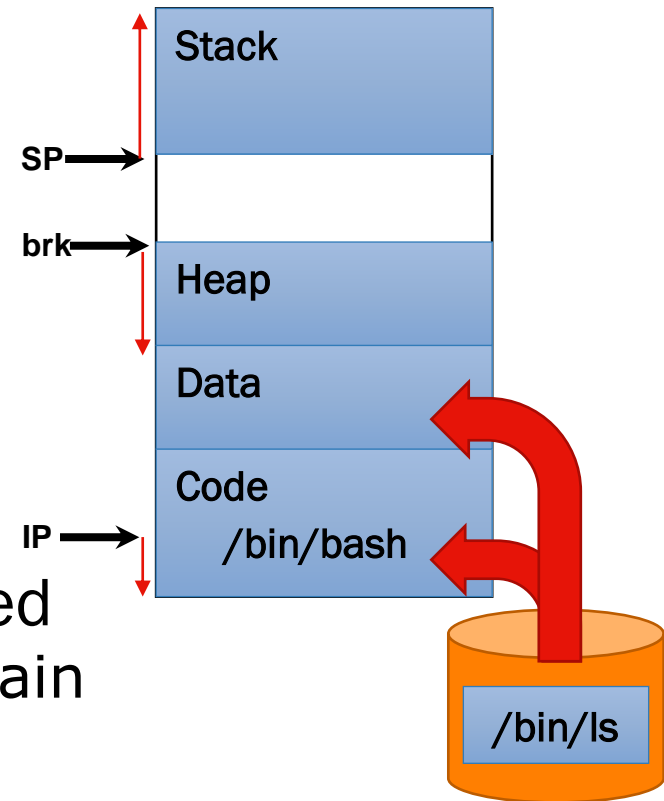
```
$ ./fork-wait  
parent 13037  
child 13039  
13037 is parent of 13039
```

```
    if (rc == 0) {                                // child  
        printf("child %d\n", (int)getpid());  
    } else {                                       // parent  
        wait(NULL);  
        printf("%d is parent of %d\n", (int)getpid(), rc);  
    }  
    return 0;  
}
```

**What is the
output without
wait()?**

exec(): Run Another Program

- Can't write entire system in one program!
 - Need to replace process with another program
- Loads program from filesystem
 - Replaces code, data, bss
 - Put argv on stack; reset %esp
 - Release heap memory
 - Reset %eip to main (really __start)
- exec() only returns to caller if failed
 - Otherwise process is now in a new main



Why Separate fork and exec?

- Lots of parameters on creating a process
 - Shell may want to
 - redirect output of children
 - change child environment
 - change child working directory
 - run child as a different user
- Hard to create simple, expressive-enough API
- Separation allows policy to be expressed in parent's program but in child's process
- Tradeoff: child may inherit things it doesn't need (or shouldn't have)

Example: Output Redirection

```
pid_t rc = fork();

if (rc == 0) {                                // child
    close(STDOUT_FILENO);                      // close fd 1
    open("./p4.output", O_CREAT|O_WRONLY|O_TRUNC, S_IRWXU); // new fd 1

    const char *myargs[3];
    myargs[0] = "wc";
    myargs[1] = "p4.c";
    myargs[2] = NULL;
    execvp(myargs[0], myargs); // runs "wc p4.c > p4.output"
} else {                                       // parent
    wait(NULL);
}
```

Termination: `exit()`, `kill()`

- When process dies, OS reclaims resources
 - Record exit status in PCB
 - Close files, sockets
 - Free memory
 - Free (nearly) all kernel structures
- Process terminates with `exit()`
- Process terminates another with `kill()`

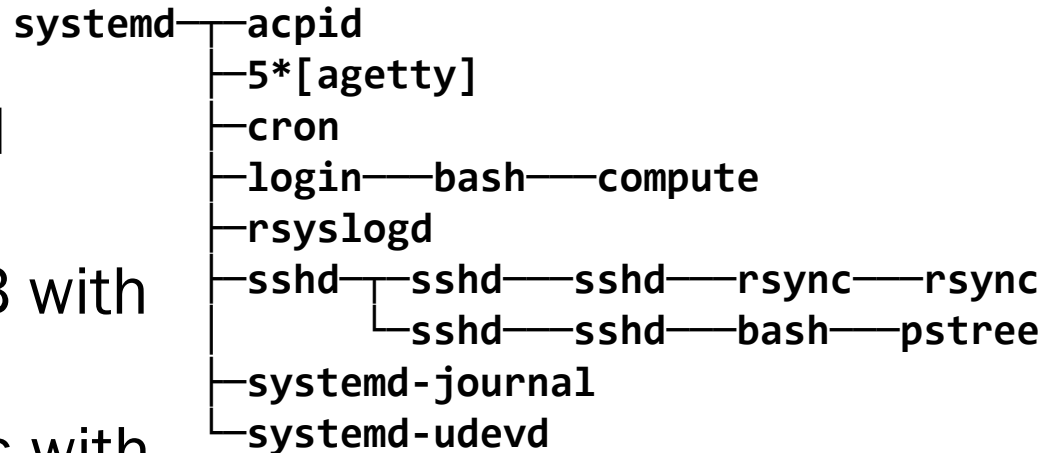
```
int main(int argc, char* argv[]) {  
    pid_t pid = fork();  
    if (pid == 0) {  
        sleep(10);  
        printf("Child exiting!\n");  
        exit(0);  
    } else {  
        sleep(5);  
        if (kill(pid, SIGKILL) != -1)  
            printf("Sent kill!\n");  
    }  
}
```

Orphans and Zombies

- Parent `wait()` on child returns status
- Must keep around PCB with status after child exit
- **Zombie**: exited process with uncollected status
- Parent exits before child?

Orphaned

- `init` adopts orphans
- Collects and discards status of reparented children after exit
- Useful for “daemons” (`nohup`)



Important Terms and Ideas

- Process, programs
- Process Control Blocks
- syscall, user/kernel mode
- Dispatcher, context switch
- Process State Machine
- New, Ready, Running, Blocked, Terminated
- fork(), wait(), exec(), exit(), kill()
- Orphans, zombies, init