

CS5460: Operating Systems

Lecture 4: Process Management

(Chapters 4, 5, 6)

Slide Credit: Andrea Arpaci-Dusseau

Assignment 1

- Due Tue Feb 2

Problem #1: Restricted Ops

How can we ensure user process can't harm others?


Solution: **privilege levels** supported by hw (status bit)

- User processes run in user mode (restricted mode) (Ring 3)
- OS runs in kernel mode (not restricted) (Ring 0)
 - Instructions for interacting with devices
 - Access to all memory
 - Ability to reconfigure CPU control registers (IDT, PTBR/CR3)

How can processes access devices?

- System calls (function call implemented by OS)
- Change privilege level through system call (trap)

System Call



`movl $6, %eax;`

`int $64`

syscall-table index

trap-table index

Kernel can access user memory to fill in user buffer
return-from-trap at end to return to Process P

Problem #2: Take CPU Away?

OS requirements for **multiprogramming**
(multitasking):

- **Policy**: Decision-maker optimizing a performance metric
 - Process **Scheduler**: Which process when?
- **Mechanism**: Low-level code that implements the decision
 - **Dispatcher** and **Context Switch**: How?

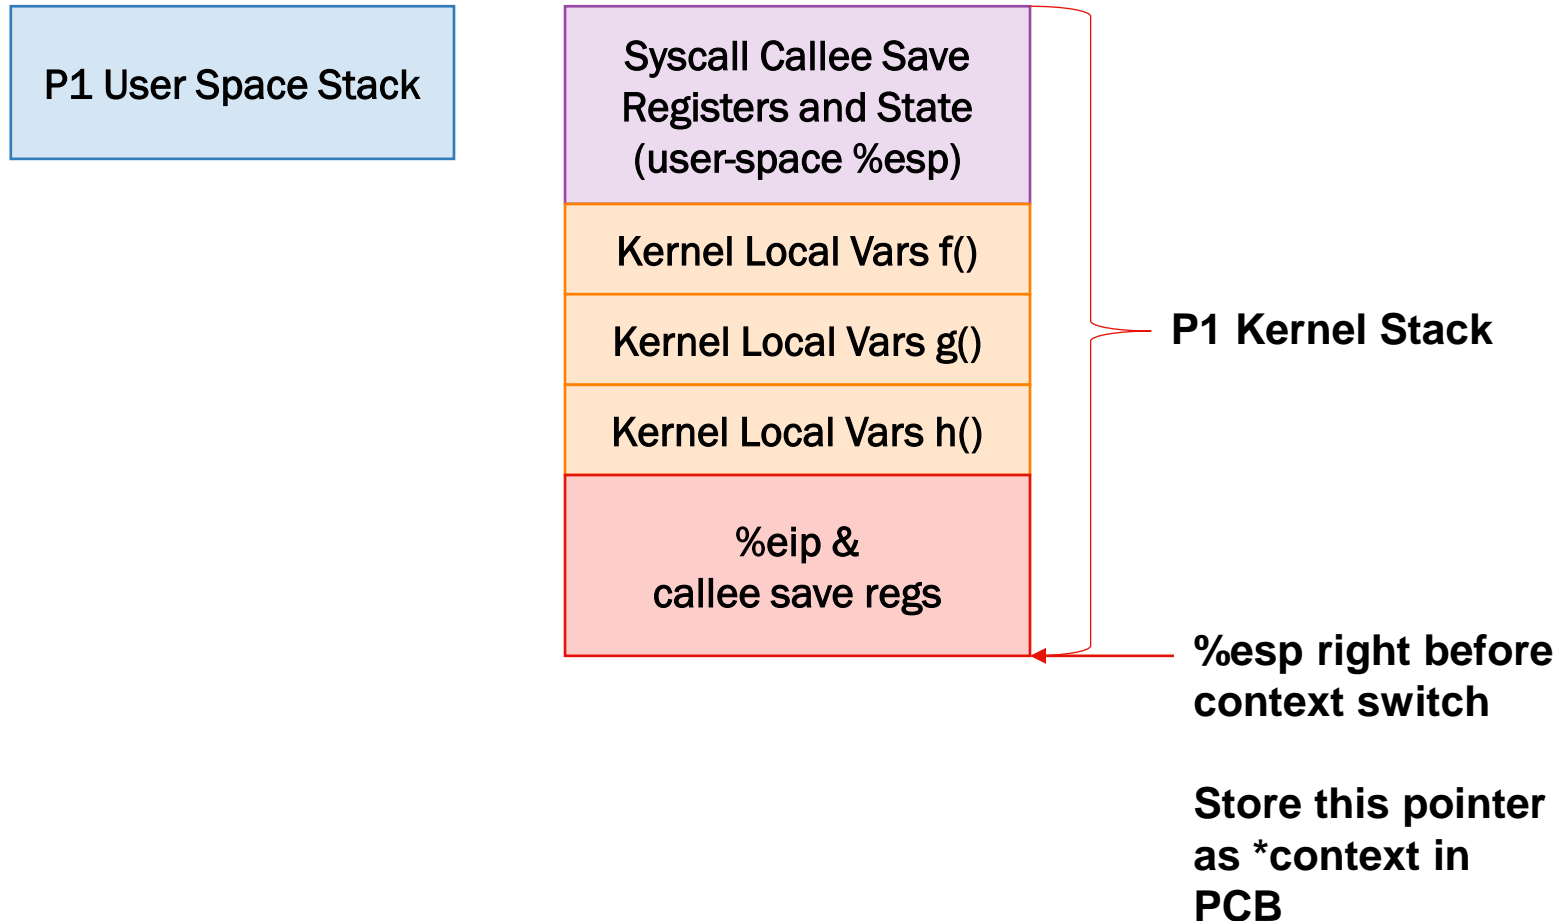
Example of separation of policy and mechanism

Q1: How does OS get control?

Option 2: Preemptive Multitasking

- Guarantee OS can obtain control periodically
- Enter OS by enabling periodic alarm clock
 - Hardware generates timer interrupt (CPU or separate chip)
 - Example: Every 10ms
- User must not be able to mask timer interrupt
- Dispatcher counts interrupts between context switches
 - Example: Waiting 20 timer ticks gives 200 ms time slice
 - Common time slices range from 4 ms to a few hundred ms

Zooming In: Registers & KStack



Operating System

Hardware

Program

Process A

...

Syscall or timer interrupt
Hw switches to kstack
Raises to kernel mode
Save regs(A) to kstack(A)
Jump to trap handler

Handle the trap
Call **swtch()** routine
Save regs(A) to PCB(A)
Restore regs(B) from PCB(B)
Switch to kstack(B)
Return to running B in kernel mode

Execute return-from-trap:
Restore regs(B) from kstack(B)
Move to user mode
Jump to B's IP

Process B

...

xv6 PCB

```
enum proc_state { UNUSED, EMBRYO, SLEEPING,
                  RUNNABLE, RUNNING, ZOMBIE };

struct proc {
    uint sz; // Proc mem size (bytes)
    pde_t* pgdir; // Page table
    char* kstack; // Bottom of kstack
    enum procstate state; // Process state
    int pid; // Process ID
    struct proc* parent; // Parent process
    struct trapframe* tf; // Trap frm for syscall
    struct context* context; // swtch() here to run
    void* chan; // If !0, sleep on chan
    int killed; // If !0, been killed
    struct file* ofile[NOFILE]; // Open files
    struct inode* cwd; // Current directory
    char name[16]; // Process name
};

struct context
{
    uint edi;
    uint esi;
    uint ebx;
    uint ebp;
    uint eip;
};
```

Context Switch

- Context switches are fairly expensive
 - Time sharing systems do 100-1000 context switches per second
 - When? Timer interrupt, packet arrives on network, disk I/O completes, user moves mouse, ...
- lab2-15 3.8 μ s
- gamow 1.6 μ s
- home 1.0 μ s
- How might one go about measuring this?

Problem #3: Slow Ops (I/O)?

On op that does not use CPU,
OS switches to other processes

OS must track process states:

Running:


On the CPU (1 on a uniprocessor)

Ready:

Waiting for the CPU


Blocked:

Asleep: Waiting for I/O or
synchronization to complete



Problem #3: Slow Ops (I/O)?

- OS maintains queues of all PCBs
 - **Ready queue**: Contains all ready processes
 - Event queue: One logical queue per event
 - e.g., disk I/O and locks
 - Contains all processes waiting for that event to complete
- Invariant: each process in 1 state and on 1 queue



CPU Virtualization Summary

- Virtualization:
Context switching gives each process impression it has its own CPU
- Direct execution makes processes fast
- Limited execution at key points to ensure OS retains control
- Hardware provides a lot of OS support
 - user vs kernel mode
 - timer interrupts
 - automatic register saving on syscall

Process Management

- OS manages processes:
 - Creates, deletes, suspends, and resumes processes
 - Schedules processes to manage CPU allocation
 - Manages inter-process communication and synchronization
 - Allocates resources to processes (and takes them away)
- Processes use OS functionality to cooperate
 - Signals, sockets, pipes, files to communicate

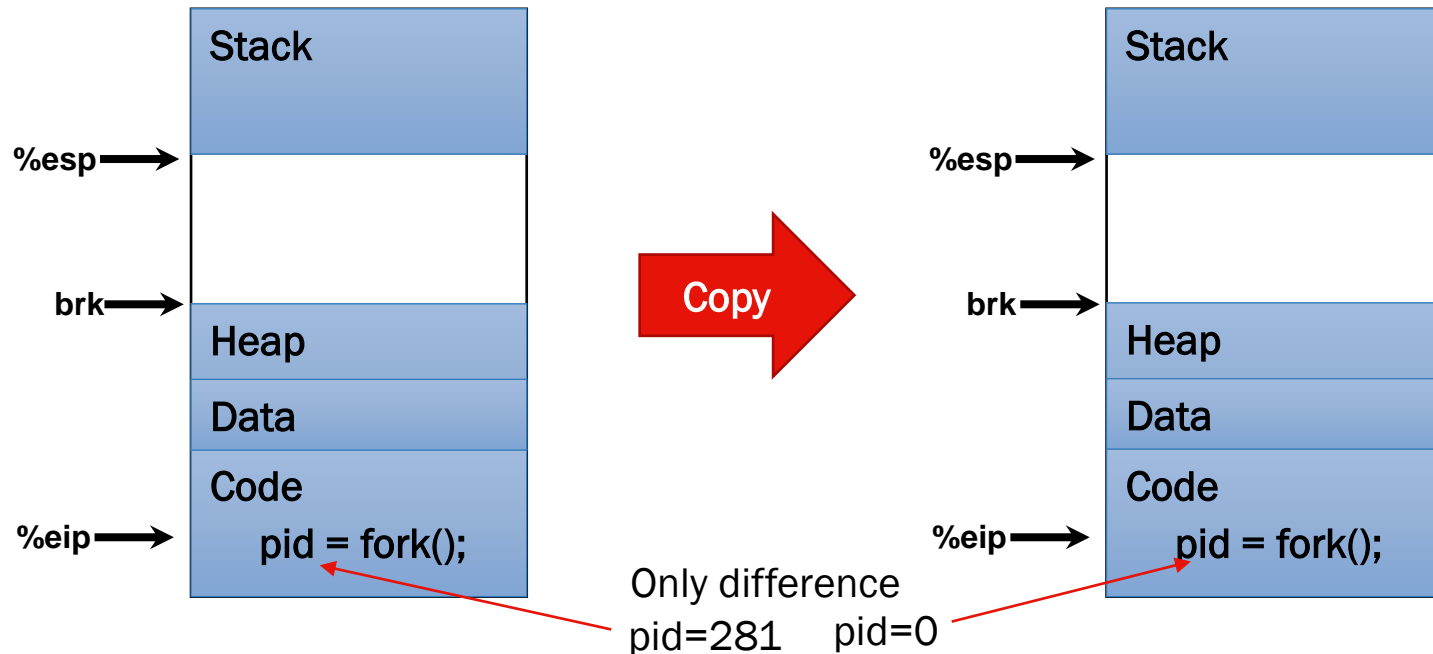
Practical Process Management

- On a Unix machine, try:
 - `ps -Af`: lots of info on all running processes
 - `kill -9 547`: terminates process with PID 547
 - `top` : displays dynamic info on top running jobs
 - Write a program that calls:
 - `getpid()`: returns current process's PID (process id)
 - `fork()`: create a new process
 - `wait()`: wait for exit of a child process
 - `exec()`: load a new program into the current process
 - `sleep()`: puts current process to sleep for specified time
- Commands work on macOS
- On Windows → Task manager (CTL-ALT-DEL)

Creating Processes: fork()

- Creates process that is near-clone of forking parent
Address space and running state is cloned
- Return of fork() differs:
0 in child
child PID in parent
- Many kernel resources are shared
open files and sockets
- wait() lets a process wait for the exit of a child
- To spawn new program, use some form of exec()

Semantics of fork()



- `fork()`, `exit()`, and `exec()` are weird!
 - `fork()` returns twice – once in each process
 - `exit()` does not return at all
 - `exec()` usually “does not return”: replaces process’ program

fork() and wait()

```
int main(int argc, char *argv) {  
    printf("parent %d\n", (int)getpid());  
    pid_t rc = fork();  
    assert(rc >= 0)
```

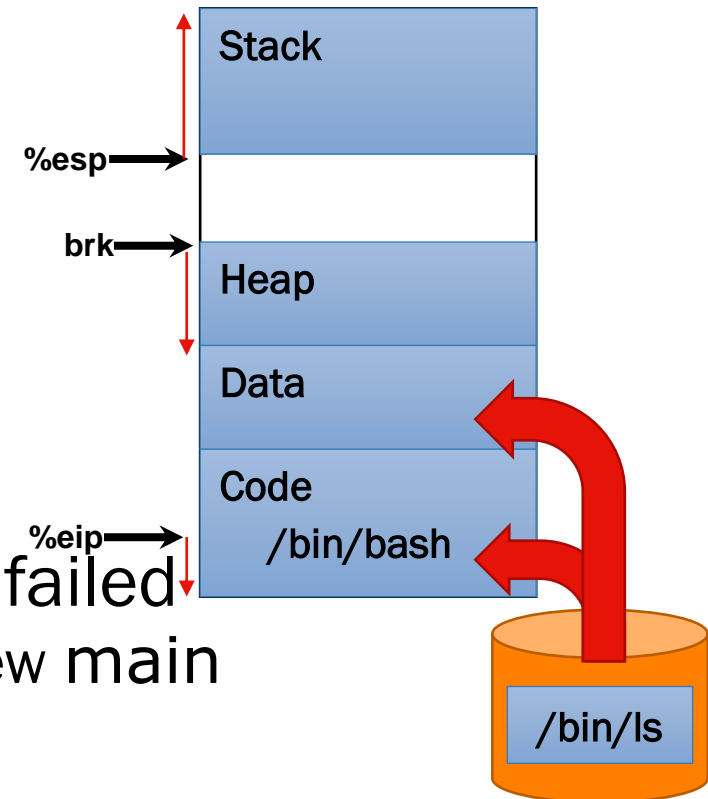
```
$ ./fork-wait  
parent 13037  
child 13039  
13037 is parent of 13039
```

```
    if (rc == 0) {                                // child  
        printf("child %d\n", (int)getpid());  
    } else {                                       // parent  
        wait(NULL);  
        printf("%d is parent of %d\n", (int)getpid(), rc);  
    }  
    return 0;  
}
```

**What is the
output without
wait()?**

exec(): Run Another Program

- Can't write entire system in one program!
 - Need to replace process with another program
- Loads program from filesystem
 - Replaces code, data segments
 - Put argv on stack; reset %esp
 - Release heap memory
 - Reset %eip to main
- exec() only returns to caller if failed
 - Otherwise process is now in a new main



Why Separate fork and exec?

- Lots of parameters on creating a process
 - Shell may want to
 - redirect output of children
 - change child environment
 - change child working directory
 - run child as a different user
- Hard to create simple, expressive-enough API
- Separation allows policy to be expressed in parent's program but in child's process
- Tradeoff: child may inherit things it doesn't need (or shouldn't have)

Example: Output Redirection

```
pid_t rc = fork();

if (rc == 0) {                                // child
    close(STDOUT_FILENO);                      // close fd 1
    open("./p4.output", O_CREAT|O_WRONLY|O_TRUNC, S_IRWXU); // new fd 1

    const char *myargs[3];
    myargs[0] = "wc";
    myargs[1] = "p4.c";
    myargs[2] = NULL;
    execvp(myargs[0], myargs); // runs "wc p4.c > p4.output"
} else {                                       // parent
    wait(NULL);
}
```

Termination: `exit()`, `kill()`

- When process dies, OS reclaims resources
 - Record exit status in PCB
 - Close files, sockets
 - Free memory
 - Free (nearly) all kernel structures
- Process terminates with `exit()`
- Process terminates another with `kill()`


```
int main(int argc, char* argv[]) {  
    pid_t pid = fork();  
    if (pid == 0) {  
        sleep(10);  
        printf("Child exiting!\n");  
        exit(0);  
    } else {  
        sleep(5);  
        if (kill(pid, SIGKILL) != -1)  
            printf("Sent kill!\n");  
    }  
}
```

Orphans and Zombies

- Parent `wait()` on child returns status
- Must keep around PCB with status after child exit
- **Zombie**: exited process with uncollected status
- Parent exits before child?

Orphaned

- `init` adopts orphans
- Collects and discards status of reparented children after exit
- Useful for “daemons” (`nohup`)



Important Terms and Ideas

- Process, programs
- Process Control Blocks
- syscall, user/kernel mode
- Dispatcher, context switch
- Process State Machine
- New (Embryo), Ready, Running, Blocked, Terminated (Zombie)
- fork(), wait(), exec(), exit(), kill()
- Orphans, zombies, init