

# CS5460: Operating Systems

## Lecture 7: Address Spaces & Address Translation

*(Chapters 13, 14, 15, 16)*

Slide Credit: Andrea Arpaci-Dusseau

# Assignment 2


- Due Tue Feb 16

# More Virtualization

- 1st part of course: Virtualization
- Virtual CPU: illusion of private CPU registers
- Virtual RAM:  
    illusion of private addresses and memory

# Motivation for Virtualization

- **Uniprogramming:** One process runs at a time



Disadvantages:

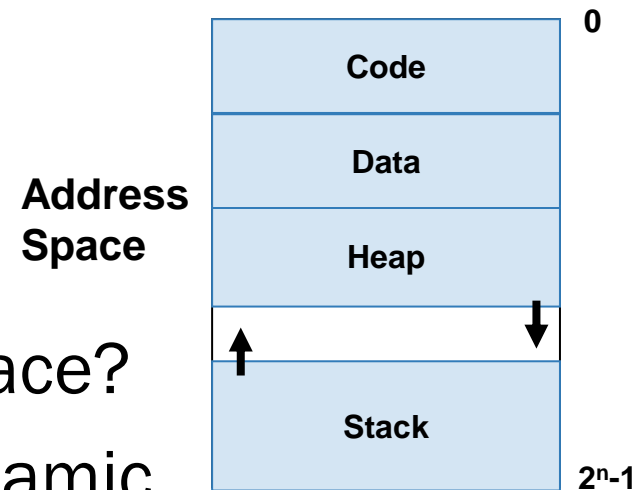
- Only one process runs at a time
- Process can destroy OS

# Multiprogramming Goals

- Transparency
  - Processes are not aware that memory is shared
  - Works regardless of number and/or location of processes
- Protection
  - Secrecy: Cannot read data of OS or other processes
  - Integrity: Cannot change data of OS or other processes
- Efficiency
  - Do not waste memory resources (minimize fragmentation)
- Sharing
  - Cooperating processes can share portions of address space

# Abstraction: Address Space

- **Address space:** process' set of addresses (that map to bytes)
  - How does OS provide illusion of private address space to each process?
- Review: What is in an address space?
- Address space has static and dynamic components
  - Static: Code and some global variables
  - Dynamic: Stack and Heap



# Motivation for Dynamic Memory

- Why do processes need dynamic memory allocation?
  - Do not know amount of memory needed at compile time
  - Must be pessimistic when allocate memory statically
  - Allocate for worst case; storage used inefficiently
- Recursive procedures
  - Do not know how many times procedure will be nested
- Complex data structures: lists and trees
  - `struct my_t *p = (struct my_t *)malloc(sizeof(struct my_t));`
- Two types of dynamic allocation
  - Stack
  - Heap

# Stack Organization

- **Stack**: memory freed in opposite order from alloc
  - `alloc(A); alloc(B);`
  - `alloc(C); free(C);`
  - `alloc(D); free(D);`
  - `free(B); free(A);`
- Simple and efficient implementation:  
Pointer separates allocated and freed space
  - Allocate: Increment pointer
  - Free: Decrement pointer
- No fragmentation
- “Automatic”: compiler adjusts stack pointer on entry/exit to calls to alloc/free space



# Where Are Stacks Used?

- OS uses stack for procedure call frames (local variables and parameters)

```
void main() {  
    int a = 0;  
    foo(a);  
    printf("a: %d\n", a);  
}
```

```
void foo(int z) {  
    int a = 2;  
    z = 5;  
    printf("a: %d z: %d\n", a, z);  
}
```

# Heap Organization

- **Heap**: memory region where alloc/free are explicit
  - Heap consists of allocated areas and free areas (holes)
- Advantage
  - Alloc lifetime is independent of call/ret
  - Works for all data structures
- Disadvantages
  - Allocation can be slow
  - End up with small chunks of free space – fragmentation
  - Leaks (forgotten free), double free
- What is OS's role in managing heap?
  - OS gives big chunks free memory to process (sbrk/mmap); library manages individual allocations (malloc/free)



# Quiz: Match that Address Location


```
int x;  
int main(int argc, char *argv[]) {  
    int y;  
    int *z = malloc(sizeof(int));  
}
```

**Possible segments: data, code, stack, heap**

Address	Location
x	
main	
y	
z	
*z	

# Memory Accesses

```
int main(int argc, char *argv[]) {  
    int x;  
    x = x + 3;  
}
```



```
0x10: movl 0x8(%rbp), %edi  
0x13: addl $0x3, %edi  
0x19: movl %edi, 0x8(%rbp)
```

**%rbp** is the base pointer:  
points to base of current stack frame

# Memory Accesses?

**Initial %rip = 0x10**

**%rbp = 0x200**



**0x10: movl 0x8(%rbp), %edi**

**0x13: addl \$0x3, %edi**

**0x19: movl %edi, 0x8(%rbp)**

**%rbp** is the base pointer:  
points to base of current stack frame

**%rip** is instruction pointer (or  
program counter)

**Fetch instruction at addr 0x10**

**Exec:**

**load from addr 0x208**

**Fetch instruction at addr 0x13**

**Exec:**

**no memory access**

**Fetch instruction at addr 0x19**

**Exec:**


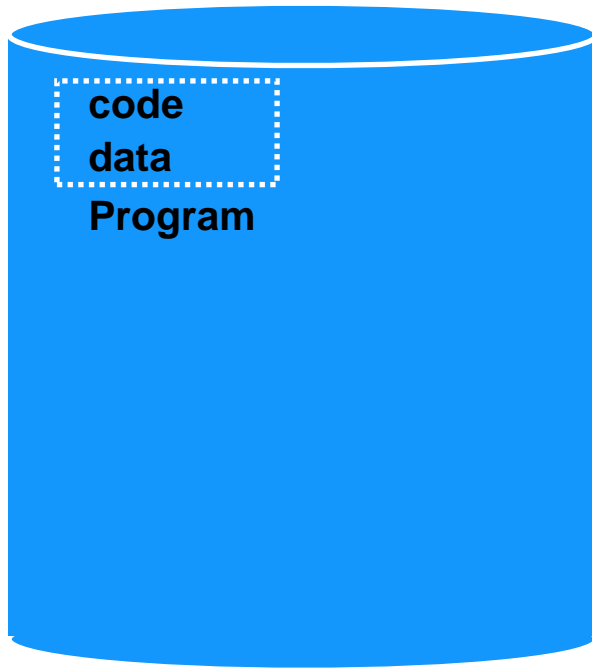
**store to addr 0x208**

# Virtualizing Memory


- How do we run multiple processes when addresses are “hardcoded” into process binaries?
- Possible solutions
  - Time Sharing
  - Static Relocation
  - Base & Bound
  - Segmentation
  - Paging (Coming Soon)


# Time Sharing Memory

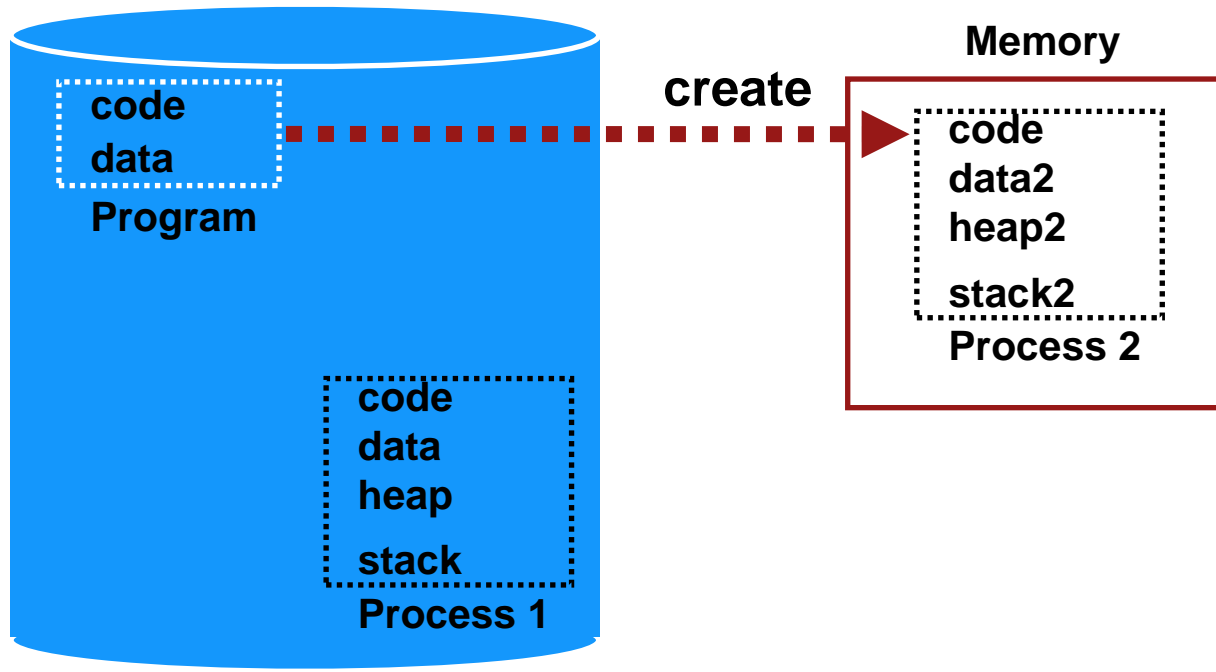
- Try similar approach to how OS virtualizes CPU
- Observation:  
OS gives illusion of many virtual CPUs by saving CPU registers to memory when a process isn't running
- Could give illusion of many virtual memories by saving memory to disk when process isn't running








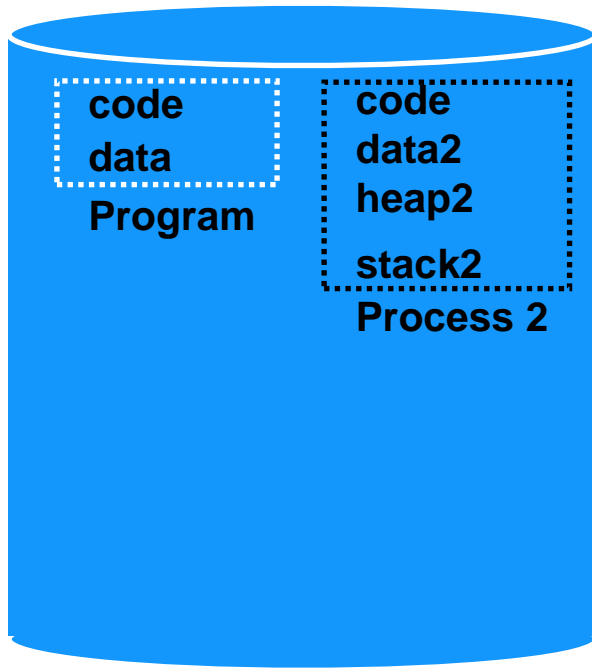












# Time Sharing is No Good

- Problem: Ridiculously poor performance
- Better Alternative: space sharing
  - At same time, space of memory is divided across processes
- Remainder of solutions all use space sharing

# Static Relocation

- Idea: OS rewrites programs before loading in memory
  - Make different process use different addresses/pointers
- Change jumps, loads of static data

```
0x10: movl 0x8(%rbp), %edi
0x13: addl $0x3, %edi
0x19: movl %edi, 0x8(%rbp)
```

rewrite



```
0x1010: movl 0x8(%rbp), %edi
0x1013: addl $0x3, %edi
0x1019: movl %edi, 0x8(%rbp)
```

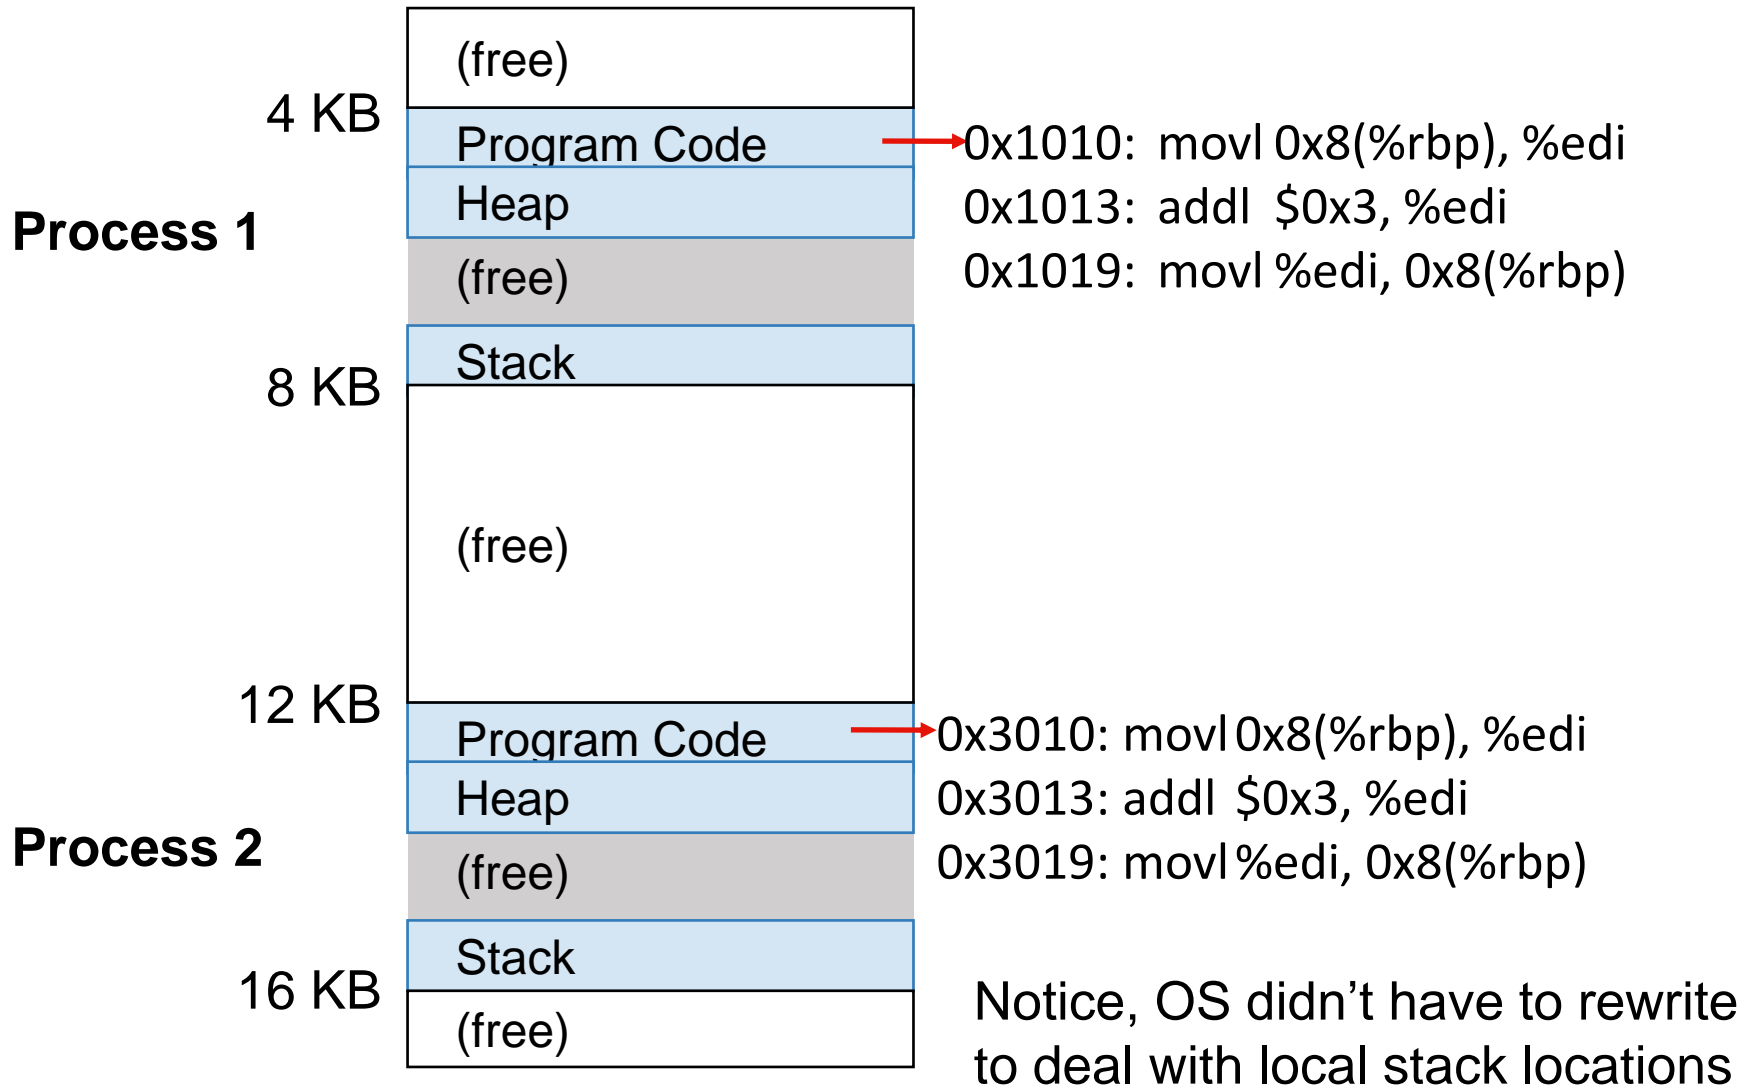
rewrite



```
0x3010: movl 0x8(%rbp), %edi
0x3013: addl $0x3, %edi
0x3019: movl %edi, 0x8(%rbp)
```



# Static: Layout in Physical Memory



# Static Relocation Problems

- No protection
  - No integrity: process can modify other processes, OS
  - No secrecy: process can access other processes, OS
- Cannot move address space after it has been placed
  - e.g. pointers in registers or on stack may refer to specific addresses, so can't easily rewrite while program running

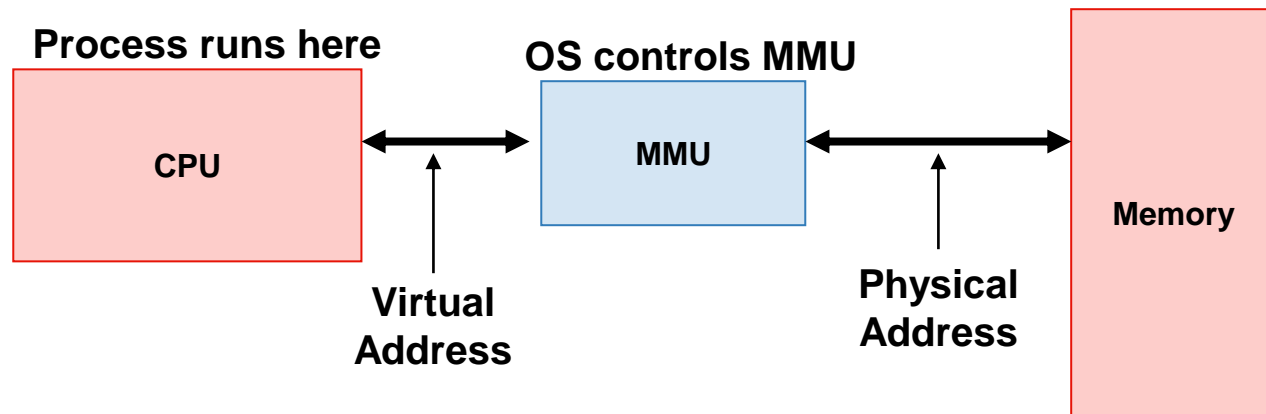
# Dynamic Relocation

Goal: Relocation at runtime and protect processes from one another

Requires hardware support: **Memory Management Unit (MMU)**

MMU dynamically changes every process address at every memory reference

- Process generates **logical** or **virtual** addresses (in their address space)
- Memory hardware uses **physical** or **real** addresses

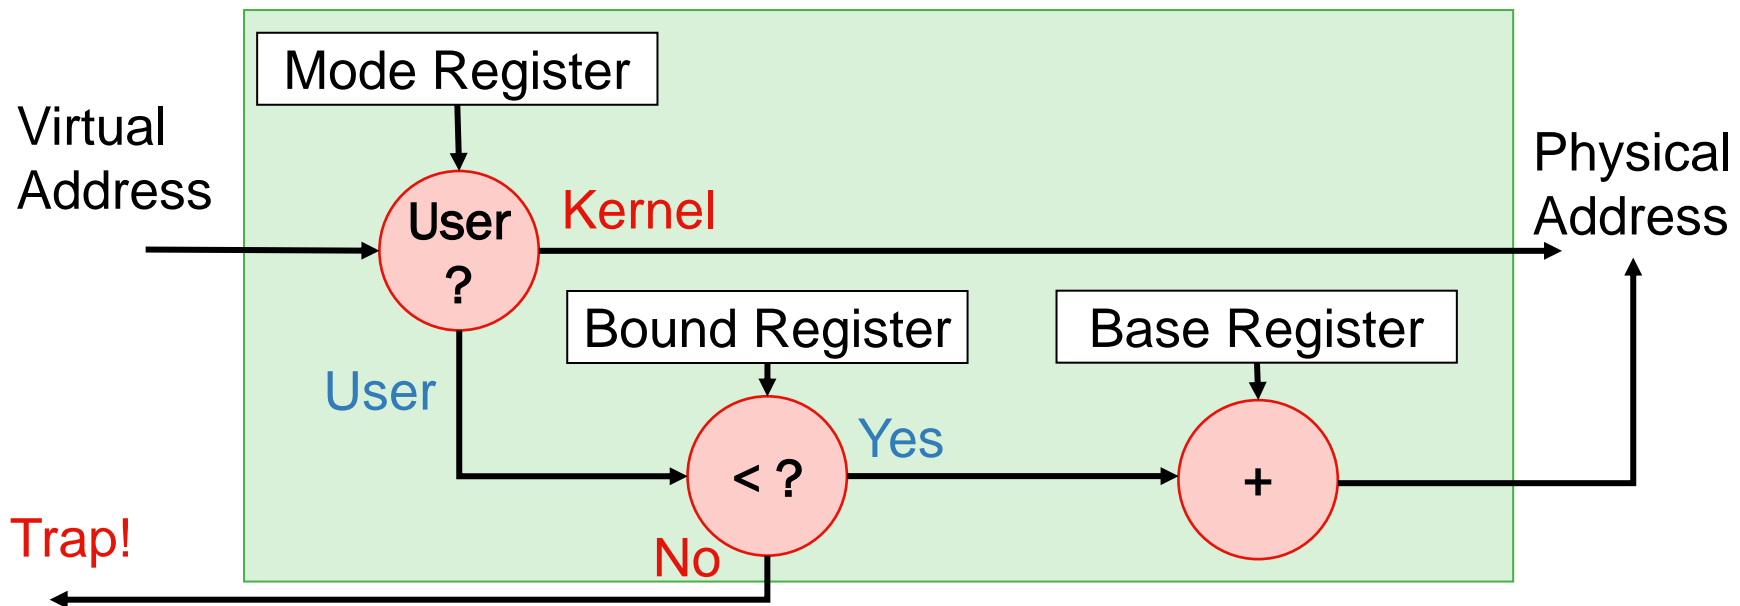


# HW for Dynamic Relocation

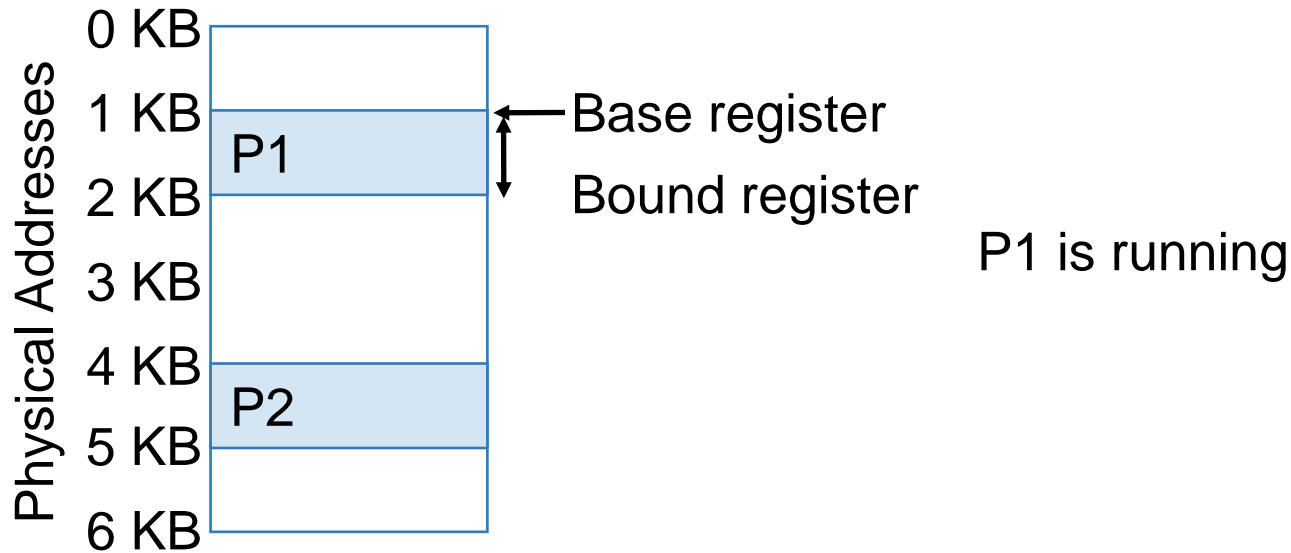
- All addresses automatically translated by MMU
  - Hardware does the translation
- OS sets up translation with privileged registers
  - Indicates where in physical memory the address space of the process starts (**base**)
  - And to what physical memory it extends (**bound**)
- User space process cannot change the registers
  - OS switches base/bound register at context switch

# Base & Bound Implementation

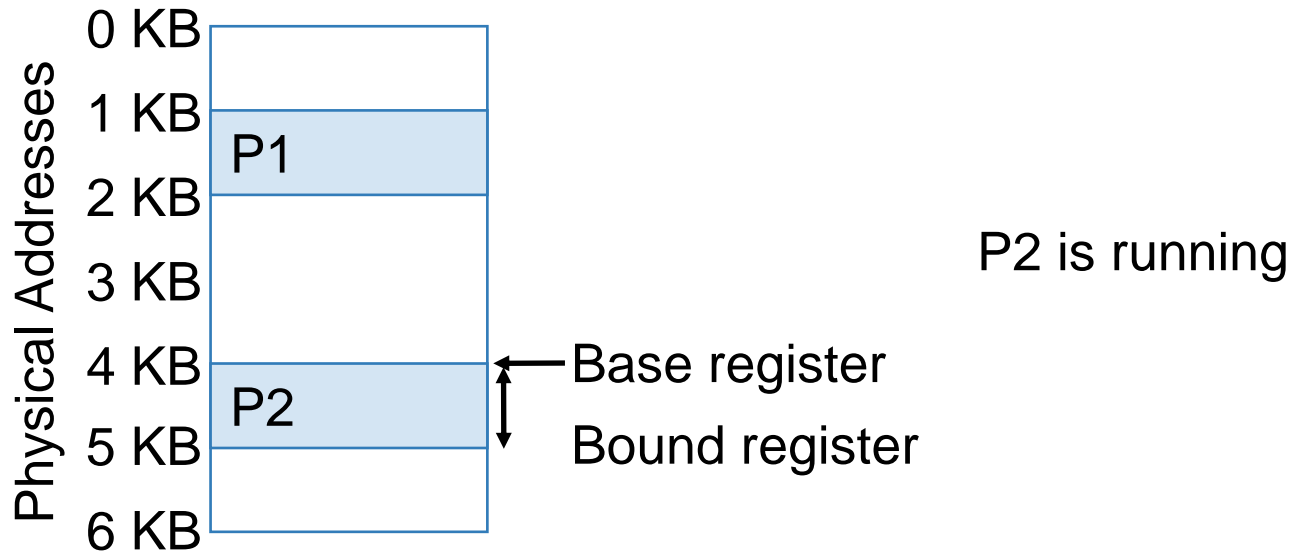
- On every memory access of user process
  - MMU compares logical address to bounds register
    - if logical address is greater, then generate error
  - MMU adds base register to logical address to form physical address



# Base & Bound Example



# Base & Bound Example



# Base & Bound Example

Physical Addresses		Virtual	Physical
	0 KB		
	1 KB	P1: load 100, %r1	load 1124, %r1
	2 KB		
	3 KB	P2: load 100, %r1	load 4196, %r1
	4 KB	P2: load 1000, %r1	load 5096, %r1
	5 KB		
	6 KB	P1: load 1000, %r1	load 2024, %r1
Can P1 hurt P2?		P1: store 3072, %r1	




# Base and Bounds Advantages

- Advantages
  - Protection (access/no access) across address spaces
  - Supports dynamic relocation
    - Can place process at different locations initially
    - Also can move address spaces
  - Simple, inexpensive implementation
    - Few registers, little logic in MMU
  - Fast
    - Add and compare in parallel

# Base and Bounds Problems

- Disadvantages
  - Each process must be allocated contiguously in physical memory
    - Must allocate memory that may not be used by process
  - No partial sharing: cannot share limited parts of address space
  - Can't control  
e.g. read/write/execute permissions



# Conclusion

- HW+OS work together to virtualize memory
  - Give illusion of private address space to each process
- Add MMU registers for base+bounds so translation is fast
  - OS not involved with every address translation, only on context switch or errors
- Dynamic relocation with segments is good building block
  - Next lecture: Solve fragmentation with paging